

---

## Overview

2+ years of experience optimizing highly parallel workloads across databases, FPGAs, and ML accelerators.

---

## Work Experience

June 2024 - **SQL Compiler Engineer**, *Snowflake*, San Mateo CA

Ongoing *Compiler Platform Team*

- Rearchitected 20+ compiler stages that are collectively run by 17M+ daily queries as part of a migration to a new expression framework.
- Engineered large-scale workload validation strategies to roll out migrated compiler stages, ensuring zero regressions across all 6B+ daily customer queries.
- Resolved 40-50 customer escalations per week during support rotations, triaging and resolving issues with query performance degradation, correctness bugs, and compilation failures across the full stack.

*Data Governance Team*

- Led design of new SQL syntax enabling multiple policies on a table, reducing boilerplate and preventing errors in applying privacy constraints.
- Designed algorithms to enforce Join Policy semantics on the parse tree of a SQL query, reducing the manual effort required to sanitize data before sharing.
- Designed optimizations on query execution plans to increase the query flexibility while maintaining privacy guarantees, reducing the amount of rewriting required for a query to satisfy constraints.

Aug 2023 - **FPGA Compiler Engineer**, *Intel*, Toronto ON

- June 2024
- Enabled the compilation of compute kernels in high-level **SYCL** code to pipelined **RTL** implementations for Intel FPGAs.
  - Created an FPGA-specific **LLVM** optimization pass in **C++** that improved performance by 15% across standard **OpenCL** benchmarks by using scalar evolution analysis to narrow induction variables.
  - Debugged complex issues across the hardware-software boundary, building expertise in the full compilation stack from high-level code to hardware implementation.

Jan 2022 - **ML Compiler Engineer (Co-op)**, *Groq*, Toronto ON

- Apr 2022
- Designed optimizations for the execution of **ONNX** models on custom hardware, enabling up to 20% throughput improvements by efficient packing of tensor operations.
  - Created optimization passes in **C++** using the **MLIR** compiler framework, applying cross-domain compilation techniques to neural network acceleration.
  - Built end-to-end validation pipeline in **PyTorch**, enabling rapid iteration between algorithmic improvements and hardware performance measurement.

---

## Education

Sep 2018 - **University of Waterloo**, *Computer Engineering B.A.Sc*, Waterloo ON

Apr 2023 Graduated with distinction.